


Identifying Code Size Optimization Opportunities in Embedded Systems

FPGA Ignite 2024

Andreas Hager-Clukas, Stefan Wallentowitz

AEMY @ Hochschule München

AEMY Lab at Hochschule München



- ▶ Research group at Hochschule München University of Applied Sciences
- ▶ Vision: Build **sA**fe **sE**cure **sM**art **sY**stems for the future
- ▶ All open source

RISC-V ISA Extensions	Open Source Chip Design Tools
WebAssembly for Embedded	Formal Methods for Embedded Systems

Open Positions @ AEMY




PhD Students

- ▶ RISC-V Instruction Set Extensions for Bytecode Virtualization
- ▶ RISC-V Instruction Set Extensions for Deeply Embedded DSP Algorithms
- ▶ Open Source Tooling for Automated Hardware-Software Co-Exploration of Fault Robustness
- ▶ Improving Inter-Tool Compatibility of Open Source EDA Tools (CIRCT, Verilator, Yosys)
- ▶ LTL Verification Support in CIRCT (LLVM framework) and Automation for ISA Extensions
- ▶ later this year: Open Source Formally Verified WebAssembly Interpreter in Rust

PostDoc

Looking for someone to help me around RISC-V instruction set extensions and open source EDA tools

Follow-up from yesterday



- ▶ Completed complex calculation unit (add and conjugation)
- ▶ cocotb test for complex calculation unit: randomization was missing
- ▶ Started implementation of the coprocessor unit: not tested
- ▶ cocotb tests started
- ▶ Now: brief review of complete project
- ▶ Find on Github: <https://github.com/hm-aemy/cvxif-tutorial>

Extending Cores

Just a very rough overview about alternatives to integrate new instructions

- ▶ ibex: The *standard* way of extending a core
- ▶ VexRiscv: Using Scala to design cores
- ▶ CVA6: Implements the core-v extension interface

Extending Cores: ibex

```
449      OPCODE_OP: begin // Register-Register ALU operation
450          rf_ren_a_o     = 1'b1;
451          rf_ren_b_o     = 1'b1;
452          rf_we          = 1'b1;
453          if ({instr[26], instr[13:12]} == {1'b1, 2'b01}) begin
454              illegal_insn = (RV32B != RV32BNone) ? 1'b0 : 1'b1; //
455          end else begin
456              unique case ({instr[31:25], instr[14:12]})
457                  // RV32I ALU operations
458                  {7'b000_0000, 3'b000},
459                  {7'b010_0000, 3'b000},
```

ibex_decoder.sv

ibex_alu.sv

```
44      //////////
45      // Adder //
46      //////////
47
48      logic      adder_op_a_shift1;
49      logic      adder_op_a_shift2;
50      logic      adder_op_a_shift3;
51      logic      adder_op_b_negate;
52      logic [32:0] adder_in_a, adder_in_b;
53      logic [31:0] adder_result;
54
55      always_comb begin
56          adder_op_a_shift1 = 1'b0;
57          adder_op_a_shift2 = 1'b0;
58          adder_op_a_shift3 = 1'b0;
59          adder_op_b_negate = 1'b0;
60          unique case (operator_i)
61              // Adder OPs
62              ALU_SUB,
63
64              // Comparator OPs
65              ALU_EQ,  ALU_NE,
66              ALU_GE,  ALU_GEU,
```

Extending Cores: VexRiscv

```
62  ✓ case class AesPlugin(encoding : MaskedLiteral = M"-----000-----0001011") extends Plugin[VexRiscv]{
63
64      object IS_AES extends Stageable(Bool)
65      object CALC extends Stageable(Bits(32 bits))
66
67      val mapping = new {
68          def DECRYPT = 25      // 0/1 =>  encrypt/decrypt
69          def LAST_ROUND = 26
70          def BYTE_SEL = 28    //Which byte should be used in RS2
71      }
72
73      //Callback to setup the plugin and ask for different services
74  ✓  override def setup(pipeline: VexRiscv): Unit = {
75          import pipeline.config._
76
77  ✓      val decoderService = pipeline.service(classOf[DecoderService])
78
```

AesPlugin.scala

Extending Cores: CVA6

```
287     parameter type x_result_t = `X_RESULT_T(CVA6Cfg, hartid_t, id_t, writeregflags_t),
288     parameter type cvxif_req_t =
289     `CVXIF_REQ_T(CVA6Cfg, x_compressed_req_t, x_issue_req_t, x_register_req_t, x_commit_t),
290     parameter type cvxif_resp_t =
291     `CVXIF_RESP_T(CVA6Cfg, x_compressed_resp_t, x_issue_resp_t, x_result_t)
292 ) (
310     output rvfi_pro
311     // CVXIF request - SUBSYSTEM
312     output cvxif_req_t cvxif_req_o,
313     // CVXIF response - SUBSYSTEM
314     input cvxif_resp_t cvxif_resp_i,
315     // noc request, can be AXI or OpenPiton - SUBSYSTEM
```

cva6.sv

But which instructions should we add?

Potential goals of instruction set extensions:

- ▶ Improve application performance
- ▶ Reduce code size
- ▶ Reduce power consumption

Computer architecture is the science of trade-offs

Key driver is an application or set of applications

- ▶ Maybe you want to focus on a single application
- ▶ In most cases you want to look at a *domain* (identified by set of applications)
- ▶ Synthetic benchmarks may not be representative

Methods to Identify new Instructions

- ▶ Goal: Identify patterns in the code that can be accelerated
- ▶ Quick evaluation of the potential benefit
- ▶ Ideally: (semi-)automated exploration

Systematic Algorithm Analysis

- ▶ Manually inspect specific algorithms (not their implementation)
- ▶ Understand how algorithm transforms into low level operations
- ▶ Identify operations in algorithms that can benefit
- ▶ Requires deep understanding of algorithms and hardware implications
- ▶ A lot of manual work, discussions; compilers are good too..

Static Analysis

Peephole optimization


- ▶ Input: Candidate patterns of instructions
- ▶ Look at small sequences of generated instructions
- ▶ Pattern matching of instruction sequences
- ▶ The more code instructions you cover, the better the optimization can be

Static Program Analysis

- ▶ How to even find candidate patterns?
- ▶ Static analysis of the code to identify patterns
- ▶ Important to define and restrict pattern search

Example Data Flow Pattern

- ▶ Each node is an operation
- ▶ Data flows between operations
- ▶ If common subpattern: candidate
- ▶ Here: $rd = rs1 \ll 4 + rs2$



Dynamic Analysis

Static analysis can find patterns in a systematic way, but if and how often instructions are executed is unknown!

Use hardware, simulation, or (often) virtual prototype

- ▶
- ▶ Dynamically executing the program exposes this information
- ▶ Need to execute on realistic data set
- ▶ Need to execute for a reasonable amount of time

Insights from dynamic analysis

- ▶ Annotate and weigh instructions with their dynamic occurrences and durations
- ▶ Potential to find patterns that are not visible statically

Turning Patterns into Instructions

Implement into hardware

- ▶ Often: Start with virtual prototype
- ▶ Ultimately: Extend core, might bring surprises

Make use in software

- ▶ Most brutal way: inline code `asm volatile (".byte 0x00, 0xC7, 0x07, 0x7B;");`
Actually this is very suboptimal, see this [stackoverflow discussion](#) for a better workable solution
- ▶ Extend compiler and assembler with intrinsics: `__riscv_complex_add(a, b);`
This needs you to provide libraries or the user to hand-craft their code
- ▶ Ideally: Compiler generates the instructions: `complex_t c = a + b;`

Adding instructions to LLVM

LLVM adds the **tablegen** language to define instructions (and more)

// This defines patterns for assembler (simple add):

```
class ALU_rr<bits<7> funct7, bits<3> funct3, string opcodestr>  
    : RVInstR<funct7, funct3, OPC_OP, (outs GPR:$rd), (ins GPR:$rs1, GPR:$rs2),  
    opcodestr, "$rd, $rs1, $rs2">;
```

```
def ADD    : ALU_rr<0b0000000, 0b000, "add">;
```

// Actual instruction selection

```
Pat<(add GPR:$rs1, GPR:$rs2), (ADD GPR:$rs1, GPR:$rs2)>;
```

With more complex instructions, more sophisticated instruction selection (ISel) must be implemented
→ makes automation hard

Generating Hardware Implementation

Ideally, you would want to also generate hardware

Commonly summarized as *Application-Specific Instruction Set Processors (ASIPs)*

- ▶ Automatically generate RTL from instruction set description
- ▶ Automatically generate compiler backend
- ▶ Commercial examples: Cudasip Codal, Synopsys ASIP Designer
- ▶ Open source example: CoreDSL

Example from our research: Sizalizer



- ▶ Goal: Reduce memory footprint of programs
- ▶ Identify further compressed instructions, in particular 48-bit
- ▶ Build a tool flow on established open source tools, each best in their domain
- ▶ Use case: DSP algorithms in the embedded space

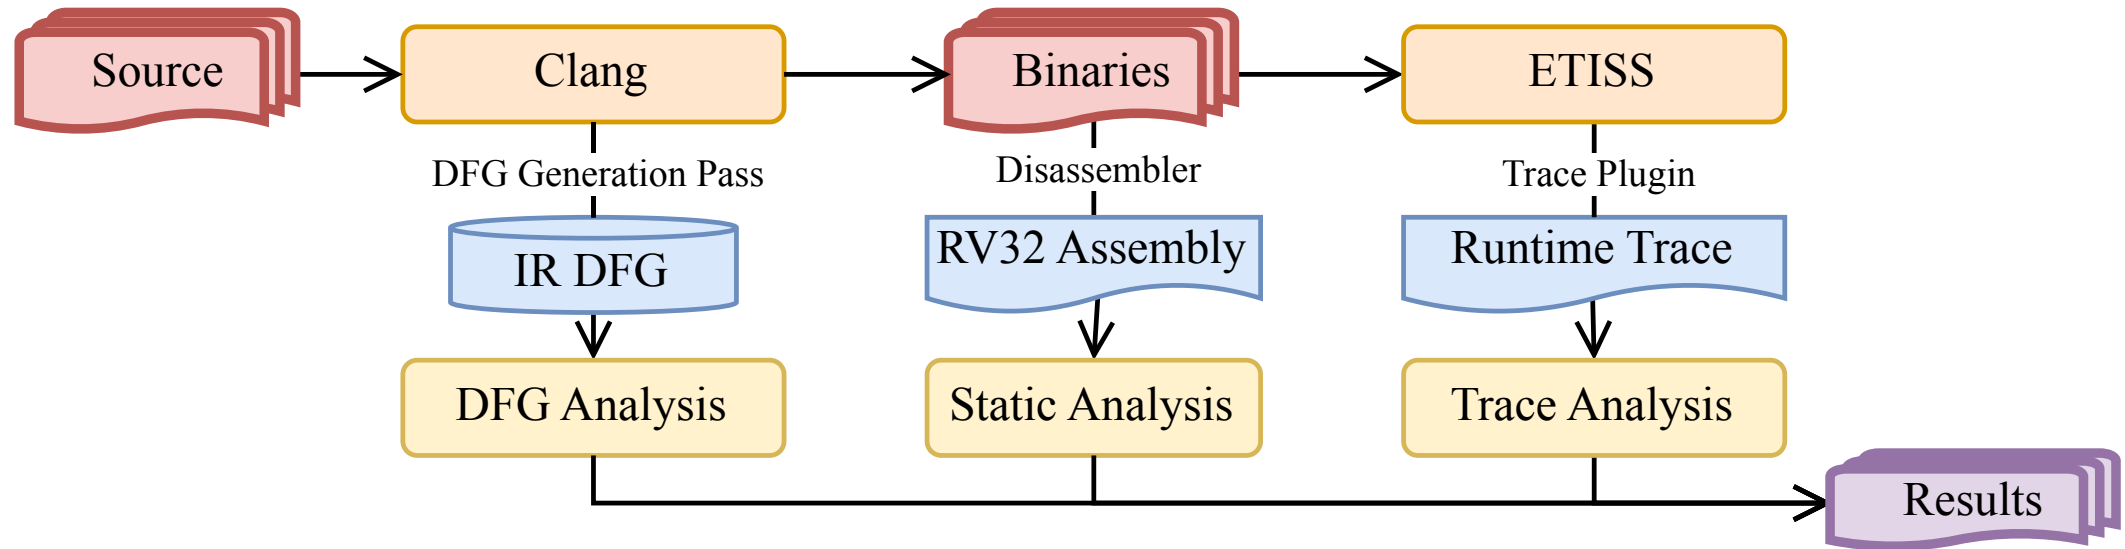
The Problem

- ▶ Tight resource constraints for embedded systems developer
- ▶ Necessity for development frameworks to achieve optimized memory footprints and power profiles
- ▶ Possible reduction of the memory footprint via an optimized instruction set architecture (ISA) with improved code density
- ▶ Revived research on domain-specific ISA optimization and the need for optimized platforms due to RISC-V
- ▶ Lack of fully integrated, easily accessible, holistic frameworks

The Vision

Sizalizer, which is an open-source framework suitable for the automated analysis of C/C++ applications targeting ISA based code size optimizations.

Analysis Framework: Architecture



Foundations

- ▶ Three levels analysis: Intermediate Representation (IR), binary executable code, and executed instructions
- ▶ Leveraging existing open-source tools to perform analysis
- ▶ LLVM IR used for structural evaluation, employing Control Flow Graph (CFG) and Data Flow Graph (DFG) metrics.
- ▶ Higher-level analysis via the Executable and Linkable Format (ELF) binary information
- ▶ Converting binary code back into assembly language mnemonics using tools like 'objdump' for evaluation

Control-Data Flow Graph


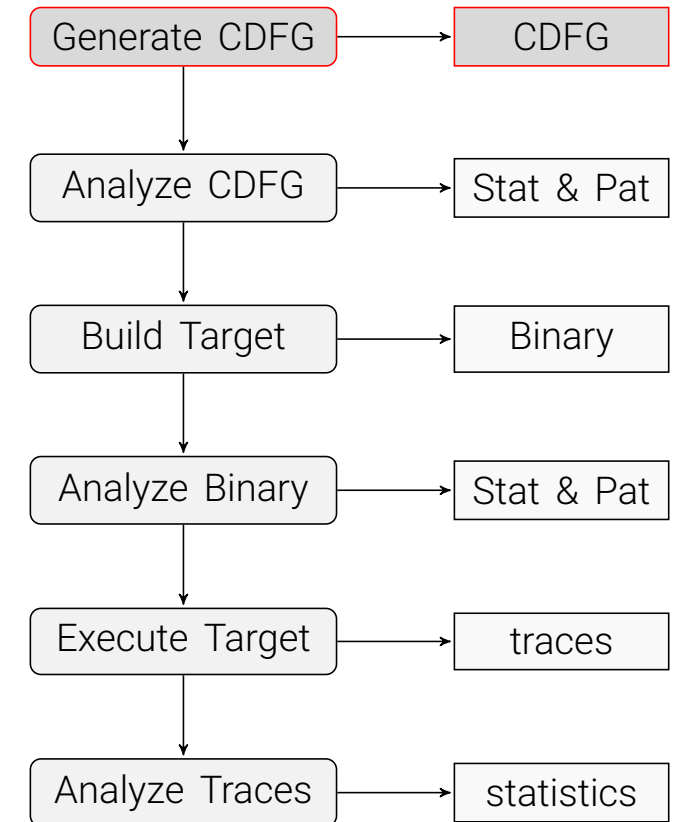


Figure: CDFG of CMSIS-DSP Example.



Control Flow Generation

```
void riscv_add_f64(  
    const float64_t * pA,  
    const float64_t * pB,  
    float64_t * p,  
    uint32_t blockSize  
) { // Start of BB % 4  
    uint32_t blkCnt;  
    blkCnt = blockSize;  
    while (blkCnt > 0U /*BB 11*/) {  
        // Start of BB % 14  
        *p++ = (*pA++) + (*pB++);  
        blkCnt--;  
    } /*BB 16*/}
```


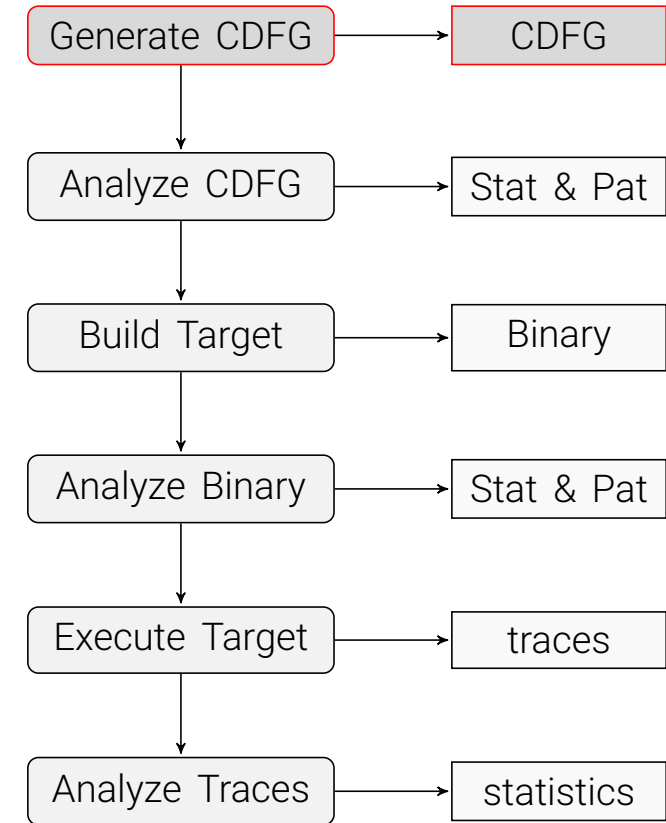


Figure: CFG of riscv_add_f64



Data Flow Generation

```
void riscv_add_f64(  
    const float64_t * pA,  
    const float64_t * pB,  
    float64_t * p,  
    uint32_t blockSize  
) { // Start of BB % 4  
    uint32_t blkCnt;  
    blkCnt = blockSize;  
    while (blkCnt > 0U /*BB 11*/) {  
        // Start of BB % 14  
        *p++ = (*pA++) + (*pB++);  
        blkCnt--;  
    } /*BB 16*/}
```


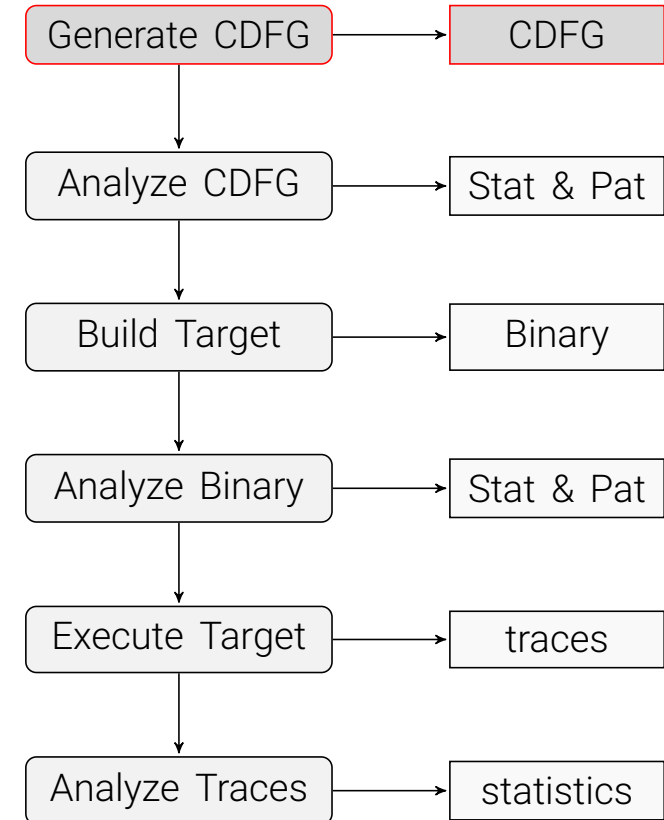


Figure: DFG of the blkCnt Variable




DFG Analysis

- ▶ Generate DFG patterns
 - ▶ N: width
 - ▶ M: depth
 - ▶ SpecialCond: func describing node relations (e.g. equality)

```
MATCH p1=(x11)-[:DFG]->...-[:DFG]-> (x1N)
MATCH ...
MATCH pM=(xM1)-[:DFG]->...-[:DFG]-> (xMN)
WHERE SpecialCond(x11, ..., x1N) AND ... AND SC(xM1, ..., xMN)
RETURN p;
```

- ▶ query_builder(depth, width, special_cond, pat_to_ignore)



DFG Analysis

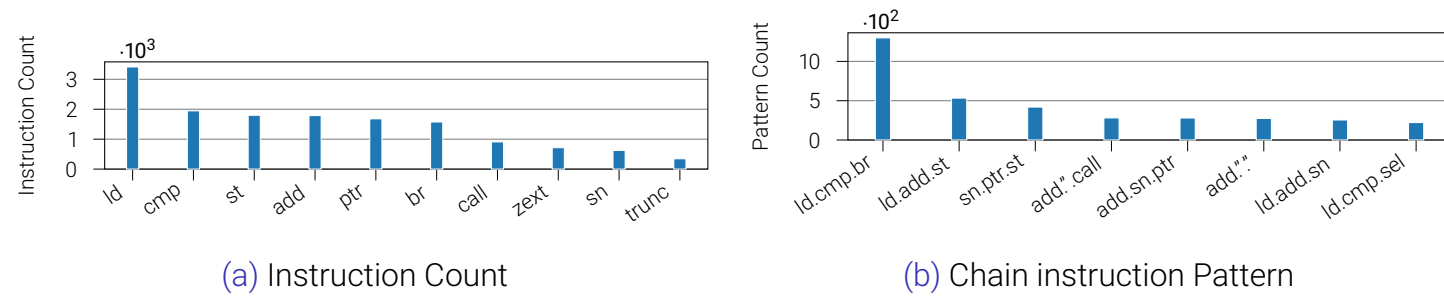
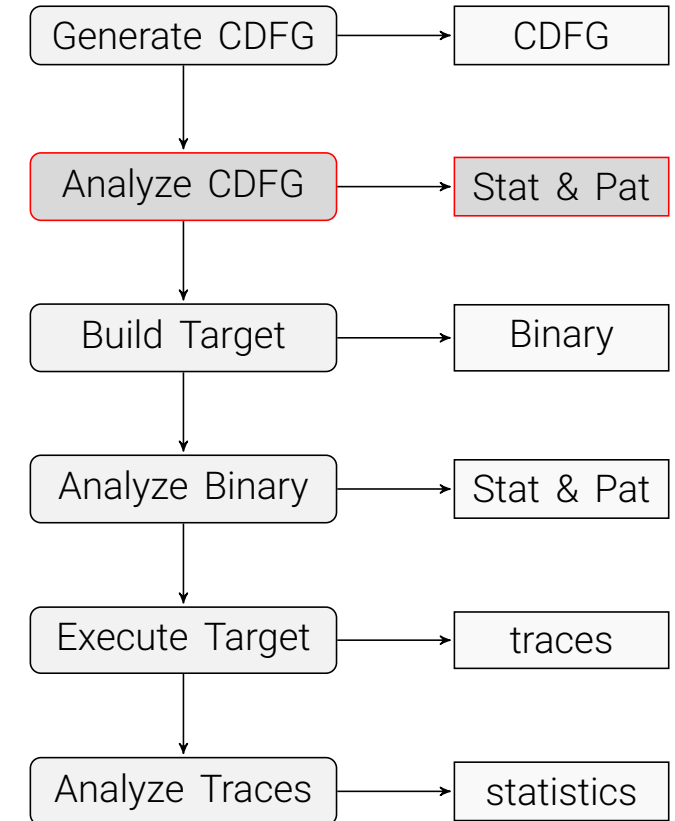


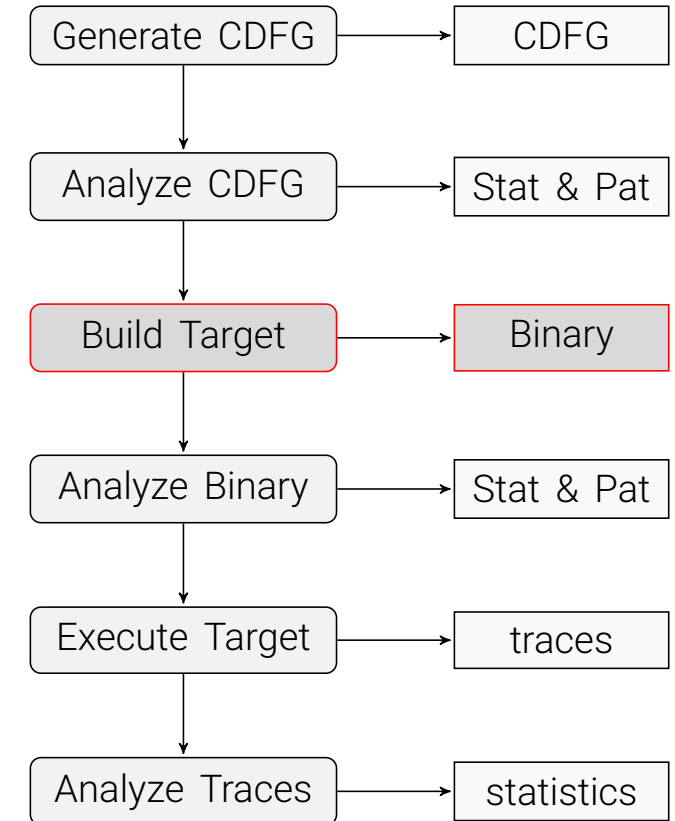
Figure: DFG Analysis Example




Binary Analysis

Address	Opcode	Mnemonic	Params
100f2	098e	c.sub	a2, a0
100f4	8145	c.li	a1, 0x0
100f6	ef003051	jal	ra, 0x10e08 ;memset;
100fa	17150000	auipc	a0, 0x1
100fe	130545e8	addi	a0, a0, -0x17c
10102	19c5	c.beqz	a0, 0x10110 ;start+0x2e;


Table: Instruction Set




Binary Analysis



(a) 16 Bit Instructions

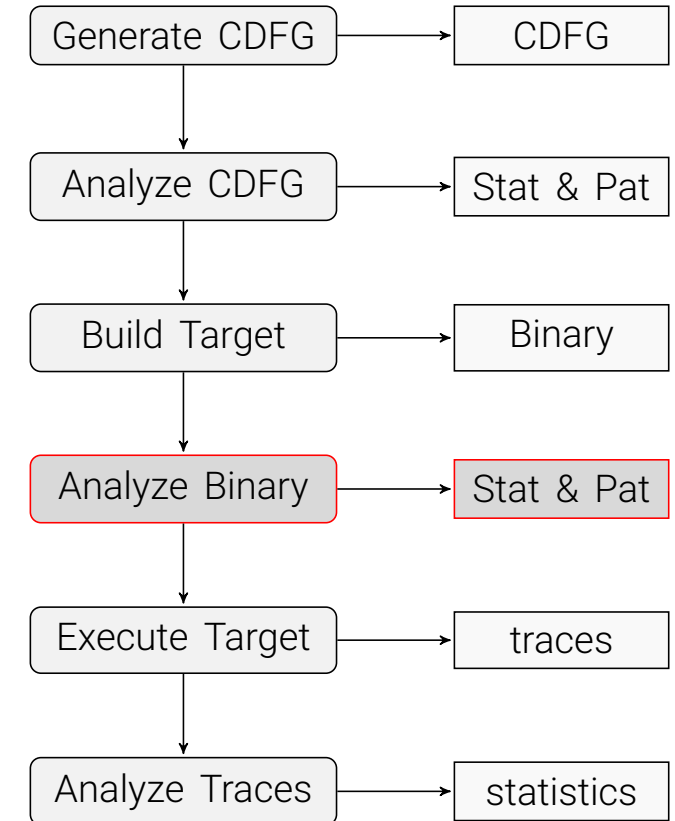


(b) 32 Bit Instructions

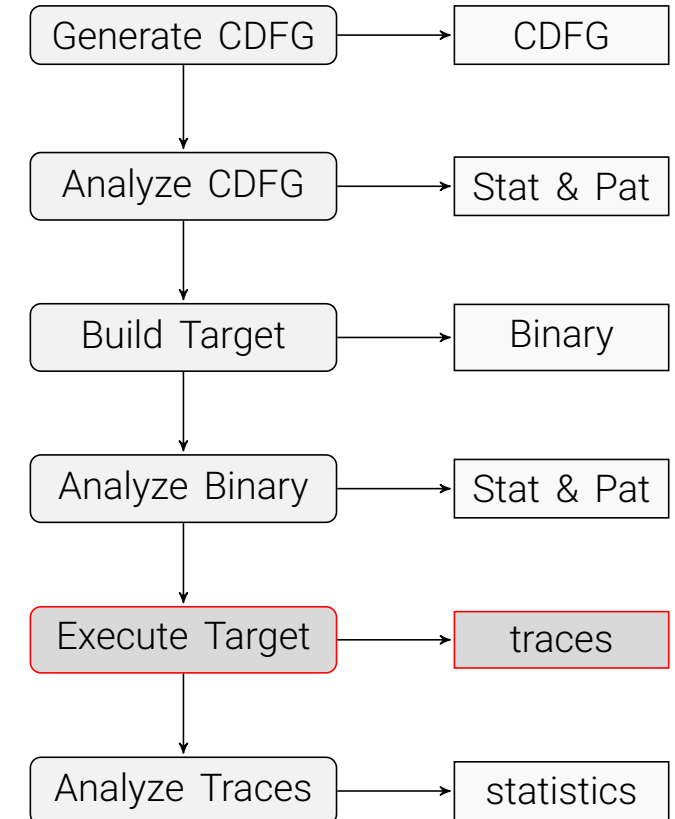


(c) Frequent instruction pairs


Figure: RISC-V Instruction Distribution Bitwith comparison.




Trace Generation

[illegible]


Analysis Framework: Trace Analysis



(a) Executed 16 Bit Instructions

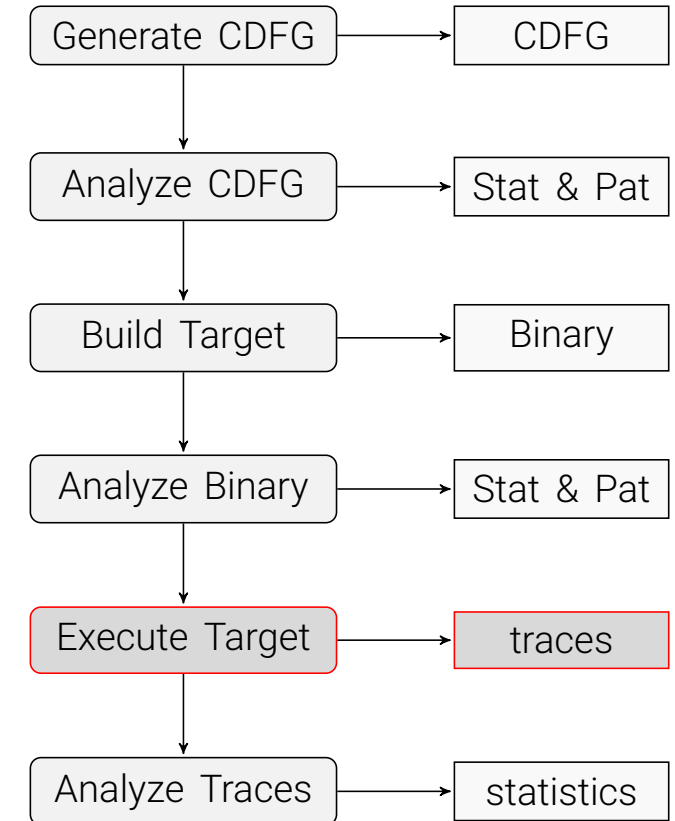


(b) Executed 32 Bit Instructions



(c) LW Immediate value Distribution

Figure: Dynamic Results



Suggested Instructions and Size Improvement




Figure: Calculated Static and Dynamic Improvement (blue: static improvement, green: dynamic improvement)

Discussion

- ▶ Sizalizer provides logical and functional instruction suggestions, but no real implementations
- ▶ Sizalizer can not provide hardware overhead of additional instructions
- ▶ Framework strenght lies in assisting the hardware architect in selecting new instructions
- ▶ Clear definition of the analysis target, beacue predefined forms or user-defined forms can be required.
- ▶ Static and dynamic size optimizations are calculated upper bounds

Using Code Analysis in Hackathon

- ▶ Sizalizer is available open source, but currently hard to use
- ▶ I have a readymade devcontainer, I can share if you are interested, but needs assistance
- ▶ Setup uses CMSIS-DSP as target for optimization
- ▶ Get in touch if this is something you would consider for Hackathon
- ▶ Alternatively: Join Jan Zielasko with his dynamic analysis tool

Contacts

Hager-Clukas, Andreas: andreas.hager-clukas@hm.edu

Wallentowitz, Stefan: stefan.wallentowitz@hm.edu

Website: <https://aemy.cs.hm.edu>